Workday Response to European Commission Consultation on
The White Paper on Artificial Intelligence - A European approach to excellence and trust
June 14, 2020

I. Introduction

Workday is pleased to provide feedback on the European Commission's White Paper on Artificial Intelligence - A European Approach (the "White Paper"). Workday is a leading provider of enterprise cloud applications for finance and human resources. Founded in 2005, Workday delivers financial management, human capital management, planning, and analytics applications designed for the world's largest companies, educational institutions, and government agencies. Organizations ranging from medium-sized businesses to Fortune 50 enterprises have selected Workday. Headquartered in Pleasanton, California, Workday has more than 12,400 employees worldwide, including 2,350 working in our 21 offices across Europe. Workday's European-based customers include: Airbus, BlablaCar, Deutsche Bank, Primark, and Siemens.

Within its applications, Workday incorporates machine learning technologies that enable its customers to make more informed decisions and accelerate operations, as well as assist workers with data-driven predictions that lead to better outcomes. Workday believes these technologies have the potential to impact enterprises in the near-term by making operations more efficient. In the longer term, enterprises will be able to reorganize operations around machine learning's unique possibilities. Promoting the thoughtful and responsible adoption of machine learning is a fundamental component of Workday's public policy agenda.

Given their impact on society, and the responsibility society has to determine how technology should be used, Workday understands the impetus for establishing a regulatory framework for machine learning. In our view, such a framework has the potential to play a constructive role in increasing the level of trust associated with these technologies. Workday urges policymakers to take a risk-based and context-specific approach to regulation that will support innovation and address social responsibility, while seeking global convergence for technologies that will cross borders. We welcome the Commission's risk-based approach outlined in the White Paper and provide the following comments that follow the White Paper's organizational structure.

II. General Comments

a. An Ecosystem of Trust

Workday is pleased with the Commission's recognition that "given the major impact that AI can have on our society and the need to build trust, it is vital that European AI is grounded in our values and fundamental rights such as human dignity and privacy protection" (p. 2). Protecting fundamental rights is a critical component in establishing "an ecosystem of trust" and should form the basis of any future AI regulation.

Indeed, trust has been a fundamental component of Workday's services since our beginning. To keep our customers' data safe and private, we have built in strong privacy protections through our privacy principles, our approach to GDPR, and our robust privacy program. We have taken a similar approach in our development of machine learning technologies. We published our Commitment to Ethical AI, which includes six key principles that guide how Workday responsibly develops machine learning (ML) for the enterprise space and works to address its broader societal impact. These principles include:

1. *Workday Puts People First.* Workday always respects fundamental human rights. Workday applies ML to deliver better business outcomes and help people in their decision-making. Workday's solutions provide customers control over how recommendations are used.
2. *Workday Cares about Society.* Workday believes that humans will always be at the center of work. Workday's focus remains on how ML can align opportunity with talent, and on contributing to the development of a machine learning-ready workforce.
3. *Workday Acts Fairly and Respects the Law.* Workday acts responsibly in its design and delivery of machine learning products and services and strives to identify, address, and mitigate bias in its ML technologies. Workday aims to ensure that machine learning recommendations are equitable.
4. *Workday is Transparent and Accountable.* Workday explains to its customers how its machine learning technologies work, the benefits they offer, and describes the data needed to power any ML solutions it offers. The company demonstrates accountability in intelligent solutions to customers and gives them a wide range of choice in how to deploy them.
5. *Workday Protects Data.* Workday's Privacy Principles apply to all of its products and services, including to its machine learning efforts. Workday minimizes the data used, and embraces good data stewardship and governance processes.
6. *Workday Delivers Enterprise-Ready ML Technologies*. Workday applies its leading quality processes – with input from customers – when developing and releasing machine learning technologies. These machine learning-powered solutions help Workday's customers tackle real world challenges.

Alongside our internal efforts, Workday is proud of our involvement in the AI ethics work championed by the EU's High-Level Expert Group on AI (HLEG). Workday piloted the HLEG's Trustworthy AI Assessment List and mapped its questions against our machine learning ethics controls. We were pleased to find high-level alignment between our internal programs and the Assessment List. By using the Assessment List to review a real Workday software tool, we were able to offer granular feedback on its effectiveness when operationalized in the business context.

b. Skills

As the White Paper notes: "The European approach to AI will need to be underpinned by a strong focus on skills to fill competence shortages" (p. 6). Workforce development and skills enhancement is one of the most critical public policy issues of our time. Already, technology—including AI and ML—is leading to ever-faster transformation of skills needed in the labor market. Fortunately, even as they transform the labor market, artificial intelligence and machine learning also provide the tools that both employers and workers need to navigate those disruptions.

As a leader in cloud applications for human resources, Workday sees broader adoption of skills-based employment practices and secure digital credentials as critical in helping define needed skills and match workers with opportunity both across the economy and across the globe. To be truly effective, skills-based practices require a common "skills language." Workday Skills Cloud uses ML to reduce 200 million+ points of skills data down to 55,000 canonical skills. Skills Cloud then helps refine job skills data, allowing employers to identify workers with the skills needed for emerging roles while also helping workers understand which emerging skills will support meaningful career advancement. When both businesses and governments use AI tools to see emerging skills gaps—in real time—they will be able to refocus on reskilling efforts dynamically.

c. International Aspects

Workday is pleased that "the EU will continue to cooperate with like-minded countries, but also with global players, on AI" (p. 8). Workday has encouraged international alignment between the United States and the EU. As others have stated, dialogue between the EU and the US could lead to the development of *de facto* global best practices for AI governance that are grounded in common democratic values. The use of intelligent technologies will cross international borders, so common rules and regulatory approaches will help spur innovation.

Workday also supports prioritizing engagement with international standards development organizations, many of which are working to address the issues raised in the White Paper. International standards have the potential to promote trust and confidence in AI. They also mitigate issues that can arise within a global marketplace when individual governments develop domestic standards that lead to unintentional barriers to trade or are used to further protectionist agendas. Workday supports the Commission's recognition of the importance of international engagement, and urges the EU to develop a regulatory regime with global interoperability in mind.

III. Specific Comments

a. Liability Regime

As our trade association, BSA, has stated, any update to EU liability rules should be guided by the following principles:

- *Liability rules should be technology neutral.* Products should not be subject to unique or heightened liability rules simply because they integrate AI and machine learning. Heightened liability rules for AI will disincentivize companies from innovating and prevent society from enjoying the many potential benefits of AI. Furthermore, given the vast array of AI use cases—from machine learning that improves human decision-making, to fully autonomous systems—a unique, "one-size-fits-all" liability regime would be difficult to implement and enforce. Instead, liability rules applicable to AI should be technology-neutral, and, just as EU product liability rules do today, allow courts to consider the risk profiles and contexts unique to various AI use cases.

- *Any changes to the EU's liability regime must be driven by a clear and demonstrated need.* The current EU product liability regime as set out in the Product Liability Directive sets out clear and time-tested rules that apply across a vast range of products, including those with embedded software. The Directive is complemented by national civil liability frameworks that reflect legal tradition and principles in each Member State, including evidentiary rules, damages regimes, and contract and tort law. Furthermore, consumers have the possibility to obtain compensation for possible harms due to AI, or other products or services under the current liability framework. Any change in liability rules without an assessment of the current regime and a demonstrated need risks chilling innovation, with little benefit for consumers.

- *Liability rules must align with the safety regime.* The EU is currently considering a new safety framework for the case of high-risk AI (e.g. to require data quality, transparency, and robustness). Any changes to the liability framework should await the conclusion of this process to ensure the regimes are aligned, as the new safety framework should inform any changes to corresponding liability rules.

b. High-Risk Determination

Workday has long recognized the need for a risk-based approach to AI and ML regulation, and thus welcomes the risk-based proposal outlined in the White Paper. Although there are general principles that should apply across all AI use cases (*e.g.,* machine learning and AI should be developed and deployed in ways that respect human dignity and rights), regulatory approaches to AI should recognize that not all AI implementations share the same impact and risk profiles. Some AI applications will pose greater risks to individuals and should be regulated as such; however, for AI applications that pose little to no risk to individuals, that same level of regulation could chill innovation and unnecessarily divert limited resources. While Workday supports the White Paper's two-pronged approach to determining "high-risk" applications of AI, the Commission should not designate specific applications of AI to be "high-risk as such" (p. 18).

Given the sheer number of types of AI applications on the market and the varying ways in which they can be deployed, understanding the sector- and use-specific implications for risk is paramount. Workday supports the

Commission's position that only a cumulative analysis of the two-pronged criteria is sufficient to classify an AI application as "high-risk." In other words, an AI application is only considered "high-risk" and subject to the commensurate legal requirements if it is *both* deployed in a high-risk sector, *and* involves a high-risk use. As the White Paper notes, the cumulative nature of this test ensures that regulation appropriately addresses high-risk AI applications without chilling innovation.

However, Workday does not support the White Paper's assertion that, in exceptional instances, certain uses of AI applications should be considered "high-risk as such," irrespective of the sector in which they are deployed (p. 18). A *per se* approach to AI regulation is unnecessary: the two-pronged risk analysis should account for *all* AI applications, including those sectors and uses that pose significantly higher risk. Risk is fundamentally a sliding scale, and the use of AI makes it no different. Absent the Commission identifying gaps in the two-pronged framework that would make it unable to account for high-risk AI, it should err against redundancy and prioritize a single risk assessment approach for all AI applications.

Aside from the fact that it is duplicative of the two-pronged analysis, reliance on *per se* risk classification could make the scope of regulation overly broad, potentially capturing low-risk AI applications that are employed in "high-risk" uses. For instance, the White Paper offers "recruitment processes" and "situations impacting workers' rights" as examples of AI uses that are high-risk *per se*. This would include those AI applications that present little or no risk to end users, but nonetheless fall under the broad category of "recruitment processes." As mentioned previously, Workday's Human Capital Management software includes a technology called Skills Cloud that offers an illustrative example of this possibility. One component of Skills Cloud is Skills Miner, a machine learning tool that collects worker-provided skills data already in Workday, like resumes and job descriptions, and classifies the data within Skills Cloud. Skills Miner yields optional skills suggestions—for candidates during their job application process as well as for hiring managers and recruiters during the job requisition creation process—highlighting skills captured by the data, including ones that otherwise would not have been represented, and providing business leaders a better understanding of skills trends, patterns, and gaps in their workforce.

c. Addressees

With respect to geographic scope of the proposed legislative intervention, Workday supports the Commission's view that legislation should apply to all relevant economic operators providing AI-enabled products or services to the EU market, regardless of whether they are established in the EU. On the question of how to distribute the legal obligations among the multiple economic operators involved in most AI supply chains, Workday takes the view that, in general, the entity that chooses to make use of the AI application should be the addressee of most legal obligations. Because the risks associated with AI are dependent on how the system is deployed, the entity that chooses to use the AI is best positioned to make judgments about whether its use conforms to relevant legal requirements. The White Paper supports this notion, recognizing that "while the developers of AI may be best placed to address risks arising from the development phase, their ability to control risks during the use phase may

be more limited" (p. 22). To the extent that certain legal obligations are better met by other suppliers, the deploying entity can shift those obligations via contract.

d. Types of Requirements

With respect to potential requirements for training data, Workday recognizes the need to ensure datasets used to train AI applications are sufficiently representative. However, requirements for training data will need to be informed by a fulsome discussion regarding the availability of public datasets. Access to sufficiently broad datasets for training represents a challenge, which can be addressed by a number of strategies, including regulatory flexibility associated with existing data sources, incentives for data partnering and sharing opportunities, the creation and public availability of government-sponsored datasets, etc. Workday encourages the Commission to couple consideration of training data requirements with consideration of strategies to ensure testing data availability.

The White Paper also contemplates the possibility of requirements for the keeping of records and data related to AI applications. While there may be instances in which it is necessary to retain records *about* the programming of the algorithm and the underlying datasets, retention of the data *itself* could be problematic from a data privacy perspective. As a human capital management provider, much of the data used in our machine learning tools is personal data, and data retention requirements may therefore conflict with GDPR requirements. Thus, Workday urges the Commission to consider alternatives to retention of the datasets themselves. In addition, in circumstances in which the AI developer is a data processor, its ability to comply with retention requirements would be dependent on the controller's agreement. These obligations should therefore be applied to the economic operator best suited to comply with them, as mentioned previously.

e. Governance

Workday supports the White Paper's notion that "the proposed governance structure should not duplicate existing functions" (p. 25). In particular, we agree that "it would be beneficial to support competent national authorities to enable them to fulfil their mandate where AI is used" (p. 24). For instance, data protection supervisory authorities could be candidates for this role, given the AI expertise they have developed by virtue of their previous work with AI systems. With respect to the proposed tasks of the governance structure, Workday supports the Commission's view that it should include a combination of responsibilities, including facilitating implementation of the legal framework through the issuance of guidance, opinions, and expertise. We have found the structure of the European Data Protection Board to be a useful model in this respect.

Workday is pleased with the Commission's acknowledgement that the governance structure must include robust stakeholder participation. Given AI's nascency, industry involvement will be crucial to ensuring that regulation is both technologically feasible and compatible with innovation. Finally, the White Paper suggests that, if the Commission adopts conformity assessments, they could be carried out by Member State-designated bodies. Workday is opposed to prior conformity assessments; such an approach would represent a significant departure

from the EU's past treatment of standalone software, which is not subject to prior authorization. However, if the Commission pursues such an approach, it should establish a mutual recognition system with other countries, recognizing that the issues at stake in Europe are truly global in nature. As suggested in the White Paper, such a mutual recognition system would ensure that an AI application is not subjected to unnecessarily redundant assessments, while advancing the Commission's broader goal of encouraging interoperability between different regulatory regimes around the world.

f. Compliance and Enforcement

Workday urges the Commission not to pursue a regulatory framework that is based on prescriptive pre-market conformity assessments for high-risk AI, which could result in barriers to market entry and chill innovation. Pre-authorization is not required for standalone software like it is for some physical products, and the fact that software includes AI does not change that equation. A potential alternative mechanism, such as self-assessments akin to the DPIAs required under GDPR for example, would be less likely to create a lag to market or unduly burden smaller enterprises. In addition, the risks that AI poses and the mechanisms for mitigating those risks are use-case and context-specific. The appropriate methods and standards for training data, record keeping, transparency, accuracy, and human oversight will vary depending on the AI use-case and context. The Commission should avoid prescriptive, one-size-fits-all, pre-market requirements, which could impede efforts to address the very risks they are intended to and significantly complicate compliance checks and enforcement.

* * *

Workday appreciates the opportunity to provide input on the European Commission's White Paper on Artificial Intelligence. We are pleased to share our support for a flexible, coordinated, and risk-based approach to AI regulation. Please do not hesitate to reach out to Natalie Mauney at natalie.mauney@workday.com for further assistance.