

Contribution to a European Agenda for AI: Improving Risk Management, Building Strong Governance, Accelerating Education and Research

A Response to :
EU Commission White Paper
“On Artificial Intelligence -- A European approach to excellence and trust”

Provided by:
The Governance of AI Research Group

Contributors include:

Adrien Abecassis, Harvard University
Justin B. Bullock, PhD, Texas A&M University
Johannes Himmelreich, PhD, Syracuse University
Valerie M. Hudson, PhD, Texas A&M University
Jack Loveridge, PhD, Harvard University
Baobao Zhang, PhD, Harvard University

Overview

As an interdisciplinary group of interested scholars and policy analysts, the Governance of AI Research Group would like to respond to the EU Consultation for the EU AI White paper. In our responses below, **we address four guiding areas where we believe the EU Commission can build upon its current framework for addressing AI**. We offer general guidance in two underdeveloped, but vitally important directions for the successful governance of AI, and we echo and expand two directions from the AI White Paper.

- (1) We build and offer recommendations for **creating a more comprehensive and holistic approach to risk management**.
- (2) We offer specific guidance on **building out governance institutions** and how these institutions need to **govern inputs to AI systems, AI systems themselves, and the uses of the AI systems** in a principles-based manner where institutions are appropriately empowered to provide oversight to this ecosystem.
- (3) We strongly support **increasing AI literacy throughout society** and give recommendations on how to do so.
- (4) We endorse **increasing funding for social sciences research** so that researchers can study the impact of AI on society to improve governance.

Recommendations for Improving the Risk-based Approach

The commission proposes a risk-based approach to guide AI regulation and outlines two cumulative criteria to determine whether an AI application is high-risk or low-risk. One of these criteria is whether an AI application is deployed in a sector that involves high risk.

We applaud approaching AI governance carefully and systematically through transparent and clear criteria as those provided by the proposed risk-assessment framework. We suggest five ways in which this risk-assessment framework can be extended or supplemented.

1. ***Going beyond the low-risk vs. high-risk binary, regulators should also consider the probability and severity of the risk, the subjects who could be impacted by the risk, and the type of risk.*** The AI white paper proposes a sector-based binary classification of AI tools as either high-risk or low-risk. We suggest three ways of enriching this risk framework to provide more detailed information.
 - a. **Probability and severity.** Risk is generally defined as the probability and severity of the occurrence of an event.¹ This definition encourages rigorous and precise modelling as a best practice. This definition of risk is moreover consistent with EU general risk assessment methodology (COM(2013)76).² The notion of “risk” in the whitepaper can hence be clarified in relation to this definition or this definition of risk (as a combination of probability and severity) could be adopted in a future version of the risk-assessment framework.
 - b. **Subjects at risk.** The risk-assessment framework could be extended by a dimension that **identifies potential or expected subjects at risk (i.e., which individuals or groups are exposed to the risk)**. The same risk will differ in its effects on different people or groups (e.g., gender, ethnicity, age group) and “subject at risk” is hence relevant information. Including a dimension of “subject at risk” in the risk assessment framework, and understanding “subject” as including collective subjects, is supported by three specific considerations.

First, building trustworthy AI systems involves building software that’s ethical and socially acceptable. Identifying the subject at a given risk helps anticipating and proactively mitigating ethical and legal problems of algorithmic discrimination and unfairness. These problems of fairness and discrimination are significant roadblocks of social acceptability that stand against the widespread

¹ Bullock, Justin B., Robert A. Greer, and Laurence J. O’Toole Jr. “Managing risks in public organizations: A conceptual foundation and research agenda.” *Perspectives on Public Management and Governance* 2.1 (2019): 75-87. <https://doi.org/10.1093/ppmgov/gvx016>.

² Available at <https://ec.europa.eu/docsroom/documents/17107/attachments/1/translations/en/renditions/pdf>

uptake of AI.³ An AI policy framework therefore should be sensitive to such social acceptability obstacles and address them by identifying subjects who are exposed to risks.

Second, identifying potential subjects of risks provides a connection to social science and participatory development. Once subjects of risks are identified, input from these affected groups and their representatives can be made to count in policymaking and AI development.⁴ Including ‘subject at risk’ as part of the risk-assessment framework builds a bridge to our recommendation concerning the importance of social science research for the systematic study of potential harm (see “Increasing Social Science Funding” below).

Third, identifying subjects of risks enables civil society activities, collective actions, and strengthens citizens’ voices. Because affected individuals and groups are identified this provides information to affected stakeholders and allows for collective actions and policy innovations on behalf of these stakeholders.

Importantly, the understanding of “subjects at risk” cannot be limited to individuals. Instead, the risk assessment framework must **take the risk for societies into account**.⁵ Other regulators do this routinely: health authorities, for example, assess the risk of an antibiotic not only based on the risk/benefit balance at the individual level, but also by integrating collective risks (strengthening the resistance of bacteria in the ecosystem, for instance). Similar mechanisms must be explicitly provided for in AI risk assessment.⁶ This is because AI tools can similarly contribute to collective harms, for example when AI decision-making is used in the information market (e.g., information bubbles, polarization, etc.).

Moreover, it has been shown that algorithms for predictive policing, loan granting decisions or college admission decisions tend over iterations to build statistical

³ Young, Matthew M., Justin B. Bullock, and Jesse D. Leczy. "Artificial Discretion as a Tool of Governance: A Framework for Understanding the Impact of Artificial Intelligence on Public Administration." *Perspectives on Public Management and Governance* 2.4 (2019): 301-313. <https://doi.org/10.1093/ppmgov/gvz014>.

⁴ For an example of participatory policymaking, see Chung, Anna, Dennis Jen, Jasmine McNealy, Pardis Emami Naeni, and Stephanie Nguyen. "Project Let's Talk Privacy." Technical Report, MIT Media Lab. 2020. <https://letstalkprivacy.media.mit.edu/ltp-full-report.pdf>

⁵ Andrew D. Selbst, Danah Boyd, Sorelle A. Friedler, Suresh Venkatasubramanian, and Janet Vertesi. 2019. Fairness and Abstraction in Sociotechnical Systems. In *Proceedings of the Conference on Fairness, Accountability, and Transparency (FAT* '19)*. Association for Computing Machinery, New York, NY, USA, 59–68. DOI:<https://doi.org/10.1145/3287560.3287598>

⁶ This should also take into account that algorithms are evolutionary throughout their life (unlike drugs whose molecules do not change), requiring specific mechanisms ensuring, for instance, that the functionalities of an algorithm affecting democracy or justice do not change without the knowledge and the consent of the people affected by this change.

groups and reinforce the differences between them, irrelative to the characteristics of the individuals composing these groups, leading to collective biases that might reinforce the polarization of societies (see also our recommendation on ‘type of risk’ below).^{7,8} Thus, **the ‘subject at risk’ should include collective risk subjects.**

- c. **Type of risk.** The risk-assessment framework could be enriched by supplementing it with information about the type of risk. Different **types of risk include health risks, privacy risks, and opportunity risks** (e.g., unfair dissemination of information about scholarships or job listings).⁹ Extending the risk-assessment framework with such an identification of type of risk invites rigor of risk analysis, offers a further bridge to social science research, and allows for connections to existing ethical and legal frameworks (such as human rights conventions).
2. **The framework should consider the upsides as well as the downsides of AI tools.** Risk is often a negative notion associated with potential harms or injuries. As any risk-based framework, the proposed guidance is purely defensive by identifying downsides. Research and practice shows that there can be various benefits to AI tools.^{10, 11,12,13} Given the existing criteria for AI that the commission has proposed, the framework could be supplemented with measures to operationalize the extent to which aspirational measures are met by an AI application (e.g. transparency, fairness, etc). Ideally, a policy to govern AI does not only include a framework to measure the downside risks but also a

⁷ Alexander D’Amour, Hansa Srinivasan, James Atwood, Pallavi Baljekar, D. Sculley, and Yoni Halpern. 2020. Fairness is not static: deeper understanding of long term fairness via simulation studies. In *Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency (FAT* ’20)*. Association for Computing Machinery, New York, NY, USA, 525–534.

DOI:<https://doi.org/10.1145/3351095.3372878>

⁸ Hadi Elzayn, Shahin Jabbari, Christopher Jung, Michael Kearns, Seth Neel, Aaron Roth, and Zachary Schutzman. 2019. Fair Algorithms for Learning in Allocation Problems. In *Proceedings of the Conference on Fairness, Accountability, and Transparency (FAT* ’19)*. Association for Computing Machinery, New York, NY, USA, 170–179. DOI:<https://doi.org/10.1145/3287560.3287571>

⁹ For an alternative risk framework see <https://ethicalos.org>.

¹⁰ Kleinberg, Jon, Himabindu Lakkaraju, Jure Leskovec, Jens Ludwig, and Sendhil Mullainathan. “Human Decisions and Machine Predictions.” *The Quarterly Journal of Economics* 133, no. 1 (February 1, 2018): 237–93. <https://doi.org/10.1093/qje/qjx032>.

¹¹ Abebe, Rediet, Solon Barocas, Jon Kleinberg, Karen Levy, Manish Raghavan, and David G. Robinson. “Roles for Computing in Social Change.” *ArXiv:1912.04883 [CS]*, January 28, 2020. <https://doi.org/10.1145/3351095.3372871>.

¹² Obermeyer, Ziad, Brian Powers, Christine Vogeli, and Sendhil Mullainathan. “Dissecting Racial Bias in an Algorithm Used to Manage the Health of Populations.” *Science* 366, no. 6464 (October 25, 2019): 447–53. <https://doi.org/10.1126/science.aax2342>.

¹³ For several practical examples of data science for social good, see: <https://www.dssgfellowship.org/projects/>

framework to assess the potential benefits.¹⁴ Such opportunities include but are not limited to: improvements to fairness, health, privacy, equity or efficiency.¹⁵

3. **The framework could assess the risk of tasks instead of sectors.** The framework proposes to assess risks based on the industry sector. We suggest that an alternative basis should be considered and we suggest “tasks” as such an alternative.¹⁶ The motivation to assess the risks of tasks instead of sectors is that sectors differ greatly internally with respect to the risk that AI tools pose. For example, although the health care sector appears to exhibit greater risks than municipal garbage collection, this need not be the case. As municipal garbage collection transitions to autonomous vehicle technology, very mundane driving decisions, such as whether the vehicles should avoid left turns, can have a significant negative impact on population safety in the aggregate.¹⁷ Likewise, accounting may as a whole appear to be a low risk sector, but individual tasks and practices that individual accountants engage in might carry significant risks. Additionally, a household appliance manufacturer that uses a simple AI classifier in customer support to score inbound requests by urgency may, in practice, discriminate by race, gender or socio-economic status. Hence, alternative frameworks should be considered that allow for a more fine-grained risk analysis than one based only on sector.
4. **The framework should be tested against general-purpose AI systems.** As AI research progresses, novel AI systems will be increasingly general-purpose. By “general purpose AI” we understand an AI system that can be deployed for more than one task. One example for that is the text generation model known as GPT-3.¹⁸ A text generation model can be used to create a customer service chatbot but also generate fake news articles.¹⁹ General-purpose AI systems hence are able to complete different kinds of tasks and be deployed in different sectors. As AI research and innovation progresses we expect that the number and proportion of general-purpose AI tools to increase. The risk-assessment framework therefore should be tested as to whether it can be meaningfully applied to assess the risks of such general-purpose systems.

¹⁴ Bullock, Justin B., Robert A. Greer, and Laurence J. O’Toole Jr. “Managing risks in public organizations: A conceptual foundation and research agenda.” *Perspectives on Public Management and Governance* 2.1 (2019): 75-87. <https://doi.org/10.1093/ppmgov/gvx016>.

¹⁵ Matthew M Young, Justin B Bullock, Jesse D Lecy, Artificial Discretion as a Tool of Governance: A Framework for Understanding the Impact of Artificial Intelligence on Public Administration, *Perspectives on Public Management and Governance*, Volume 2, Issue 4, December 2019, Pages 301–313, <https://doi.org/10.1093/ppmgov/gvz014>

¹⁶ Bullock, Justin B. “Artificial intelligence, discretion, and bureaucracy.” *The American Review of Public Administration* 49.7 (2019): 751-761. <https://doi.org/10.1177/0275074019856123>

¹⁷ Himmelreich, Johannes. “Never Mind the Trolley: The Ethics of Autonomous Vehicles in Mundane Situations.” *Ethical Theory and Moral Practice* 21, no. 3 (May 17, 2018): 669–684. <https://doi.org/10.1007/s10677-018-9896-4>.

¹⁸ Brown, Tom B., Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, et al. “Language Models Are Few-Shot Learners.” *ArXiv:2005.14165 [Cs]*, June 4, 2020. <http://arxiv.org/abs/2005.14165>.

¹⁹ Kreps, Sarah and Miles McCain. “Not Your Father’s Bots: AI Is Making Fake News Look Real.” *Foreign Affairs*. August 2, 2019. <https://www.foreignaffairs.com/articles/2019-08-02/not-your-fathers-bots>

5. ***Risk should be anticipated already in the development phase.*** The framework treats risk as a given quantity that is classified throughout the framework. As such, the framework is intended to apply to an AI application at the point of deployment. We suggest that the importance of responsible innovation should be emphasised as part of a risk-assessment framework. The probability and severity of risk that eventually results from an AI application can be anticipated and mitigated during the design and development of the AI application. The risk of an AI application is, at least to some limited extent, under the control of those who develop the AI application. This aspect of technological development and the importance of responsible innovation could be recognized more prominently in AI regulation and, perhaps, even be reflected in a possible extension of the risk-based framework consistent with other regulations and guidelines.²⁰

²⁰<http://www.businessofgovernment.org/sites/default/files/Risk%20Management%20in%20the%20AI%20Era.pdf>

Guidance on Building Governance Institutions

We would like to begin this section with the reminder that the discussion on the different levels of regulation and governance could be clarified in the paper. Algorithm regulation and governance could be carried out at **three levels** of the algorithmic process, from downstream to upstream:

Regulating the social impact of algorithms downstream, setting rules and principles of what they should or should not produce, up to developers to elaborate the technical solutions to comply with them. This is the approach of many international declarations on AI (cf. the OECD principles: AI must be "human-centered" and comply with the imperatives of "inclusive growth, sustainable development, well-being, equity, transparency and explicability, robustness, safety"...), often raising questions of enforcement.

Regulating the algorithms themselves to prevent harmful effects: auditability, certifications, in some cases full transparency of the source code and data used, etc. This is an emerging discussion, to which the white paper devotes broad aspects (and to which companies often resist the most).

Regulating upstream what feeds the algorithms, i.e. the data: this involves setting rules and principles on who has access to what, who has the right to use what, under what conditions, with what consent or authorization. This is the level at which the GDPR applies, as well as many of the discussions on privacy, and the Commission's second White Paper on the European Data Strategy whose link with this White Paper should be clarified.

Each level is necessary but by itself incomplete and effective regulation will need a combination of these different levels with varying weights depending on the domains, applications, and type and level of risk. The Commission should clarify the levels which it wishes to prioritize and how this relates to other existing regulations and international declarations and commitments to which the EU has subscribed.

With these levels of regulation and governance in mind, what follows are some "principles of enforcement" concerning AI applications, especially when used as decision making systems, plus thoughts on mechanisms permitting that enforcement:

1. Principles of Enforcement

- a. Everyone has the right to **know** when they are engaging with an AI system. They should be notified and be shown a standardized identification label with the contact information for the liable party (see below).²¹

²¹ See, for example, Brian Higgins, "Recognizing Individual Rights: A Step Toward Regulating Artificial Intelligence Technologies," *Artificial Intelligence Technology and the Law*, 10 January 2018, <http://aitechnologylaw.com/2018/01/recognizing-individual-rights-regulating-ai/>

- b. Everyone has the right of **appeal** to a human being in the decision making entity who is empowered to make a decision without recourse to AI applications. Only a human being can see when an algorithm has veered from its intended purpose.²²
 - c. Everyone has the right to **litigate** harm caused by AI applications. The liability rests with the vendor of the algorithm who has sold the product to the entity whose deployment of the AI system caused harm. In some cases, that may be the same party as the vendor, but often it is not. End party entities that were negligent in their purchase from the vendor, having failed to check for due diligence on the part of the vendor with regard to appropriate auditing, may also be held partially liable.²³
2. **Harm Avoidance Mechanisms:** The rights to **know**, to **appeal**, and to **litigate** will require mechanisms to make their implementation possible:
- a. Deployment of AI applications will require capabilities to ensure:
 - i. Mandatory notification given to all who encounter the system, accompanied by
 - ii. Mandatory labeling identifying who was the vendor who sold the system to the decision making entity deploying the system. In some cases, the vendor and the entity are the same.
 - b. Within each decision making entity that utilizes AI applications, a Human Appeals Department or its like should be established. Those who staff this department must understand the intentions of the decision making process. In addition to being enabled to overturn or revise AI decisions, this department must also be empowered to demand changes to the algorithms be made by the vendors where the algorithms have been identified as veering from the original intent of the customer.
 - c. Due diligence on the part of customers of these vendors of AI applications demands that there be a certification process for algorithms before they are sold or deployed, whereby independent external auditors can certify that to the best of their knowledge, a particular algorithm does not cause harm. "Harm" in this case is defined not only at the individual level, but also at the collective level, and pertains to tangible harm such as job loss or physical injury, as well as less tangible harm, such as the reification of bias against protected classes within a society. All source code, as well as testing data, logs, and other pertinent material, must be made available to independent external auditors for such certification. Changes to the algorithm over time must also be submitted for certification. Should an entity purchase a system without such certification, the entity would have demonstrated negligence and thus would become equally liable with the vendor for any harm caused.

²² See, for example, Sherif Elsayed-Ali, "Why Embracing Human Rights Will Ensure AI Works for All," World Economic Forum, 13 April 2018,

<https://www.weforum.org/agenda/2018/04/why-embracing-human-rights-will-ensure-AI-works-for-all/>

²³ An excellent resource is Rashida Richardson, Jason M. Schultz, and Vincent M. Sutherland, "Litigating Algorithms," AI Now Institute, September 2019, <https://ainowinstitute.org/litigatingalgorithms-2019-us.pdf>

- d. In litigation, plaintiffs asserting an AI system has caused them harm are entitled to access the source code and any other relevant material of the algorithm under question, and vendors must waive any rights to secrecy or opacity regarding the code.
- e. The government must have the power to prevent or reverse the deployment of a system where it cannot be shown that risks have been satisfactorily mitigated.
- f. Retirement of an AI tools system must ensure that claims can continue to be litigated. This means that source code, testing data, etc., must still be retained by the vendor even if the vendors retire the system. Retention may also be accomplished through the certification companies or independent archives established for that purpose, as well.

3. **What Would Be Required to Establish These Mechanisms**

- a. The first need is to stand up a cadre of trained independent algorithm auditors. This will require an accelerated educational effort, and the establishment of independent companies whose certifications will carry weight because they are both impartial and independent.
- b. The second need is to develop the capabilities of lawyers and legal specialists to litigate algorithm liability cases. This, too, will require an educational effort by law schools.
- c. A new track within the Human Resources field will need to be created that focuses on the training of new HR Human Appeals Officers for AI tools.
- d. Regulations governing the new mechanisms will require the creation of public institutions such as agencies, departments, and commissions, as well as the enumeration of their powers to prevent or reverse deployment of AI tools.
- e. Education of the public concerning their rights to know, to appeal, and to litigate must be promulgated, and citizens should ideally understand in broad terms how AI systems function.

We provide further detail concerning these requirements in the following section.

Increase AI Literacy Skills Within the Public Sector, the Private Sector, and the General Public to Enable Robust Auditing and Appeals Capability

To successfully regulate AI applications, the EU will need to invest significantly in human capital by training government employees, particularly regulators, as stated on page 6 of the white paper, training private sector employees and developers, and educating the general public. Building AI literacy is essential for those making decisions about the funding, procurement, and deployment of AI systems.²⁴ Algorithm auditors are needed to evaluate the impacts of machine learning algorithms; they should be competent in addressing the following questions:

- What are the costs and benefits of deploying the AI application? In particular, what is the level of risk to individuals and groups if a given AI application is deployed?
- How has the AI application been tested for fairness, robustness, and safety, not only with regard to individuals, but also with regard to groups and collectives?
- What risks can be predicted to arise when the AI application interacts with other systems? For example, credit score AI systems may interact with unintended negative effects with the loan application AI systems of particular banks.

The AI White Paper encourages adoption of AI applications by the public sector. Therefore, public sector employees should be trained to evaluate AI applications before procurement. Those running public agencies need to understand the various risk and technical aspects associated with use of AI applications and systems. Managers within these agencies will also need to understand how to design information flow systems with AI applications in mind, they too will need to understand the risk based decision making strategies illuminated here and throughout the literature. Even so, the government will not have the resources to do the job alone.

Therefore, we identify three specific high need areas for human capital development; these include:

1. Recruit and train independent algorithm auditors. The auditors should be trained by entities independent of vendor tech companies, such as universities. We envision that for a healthy AI ecosystem to exist, independent and impartial auditing firms that check the work of developing vendor auditors will be necessary in the private sector.²⁵ The government will not have the resources to perform all the auditing checks required. But

²⁴ Horowitz, Michael and Lauren Kahn. "The AI Literacy Gap Hobbling American Officialdom." *War on the Rocks*. 14 January, 2020, <https://warontherocks.com/2020/01/the-ai-literacy-gap-hobbling-american-officialdom/>

²⁵ Clark, Jack, and Gillian K. Hadfield. "Regulatory Markets for AI Safety." arXiv preprint arXiv:2001.00078 (2019).

insofar as the government must have its own capability in auditing, as well, these accreditation programs should be made available to public sector employees.

2. Develop the capabilities of lawyers and legal professionals to litigate algorithm liability cases. Law schools should offer courses and training in AI and the law. Lawyers and legal professionals will also benefit from basic education in the auditing of algorithms, since many expert witnesses in these cases will be professional algorithmic auditors.
3. Create a new track within the Human Resources field that focuses on the training of new HR Human Appeals Officers for deployed AI systems. The HR field already trains professionals in relevant skills, such as grievance mediation, identification of discriminatory differential treatment, legally compliant decision making with regard to hiring and other direct effect processes, among many others. Training HR professionals to handle legally mandated human appeals processes for AI systems is a natural extension of this field of expertise.

Finally, the general public, the end users of so many of these AI applications, need to be better educated on the basics of how AI systems work and how they impact their lives. A widespread Public Service Announcement (PSA) campaign both in physical and digital space is needed to raise awareness of the inputs to AI systems, to forms of learning algorithms, and how individuals and society are impacted by their interactions with these systems. These PSAs could be delivered in partnership with other like-minded institutions, and should emphasize the public's right to know, to appeal, and to litigate. Members of the general public must never feel helpless or ignorant in the face of AI system deployments, which we predict will become increasingly pervasive over time.

Increase Funding for Social Science Research

One key element in training and developing the needed human capital for society is increasing funding for social science research. Social science research can inform policymakers and citizens on how to use AI applications in ways beneficial to society. Additionally, social scientists could counterbalance the influences of AI firms by identifying the flaws in AI systems. To this end we provide the following comments:

1. **Social science research has uncovered major flaws in automated systems.** Information studies scholar Safiya Noble uncovered racial and gender bias in search engine algorithms.²⁶ Political scientist Virginia Eubanks documented how errors in algorithms used in welfare decision-making, housing allocation for homeless people, and preventative child protection interventions further harm marginalized communities.²⁷ Economist Sendhil Mullainathan, working with public health researchers, detected anti-Black racial bias in a widely-used health care algorithm.²⁸ In their social science adjacent work, computer scientists Joy Buolamwini and Timnit Gebru found that commercial facial recognition systems are biased against Black people, particularly Black women.²⁹ These examples show how central social science research is to exposing serious harms that could be caused by AI systems.
2. **The gap between research in AI and relevant work across the social sciences is growing and risks obstructing opportunities for productive collaboration.** Recent work published in *Nature Machine Intelligence* resulting from research led by Dashun Wang at Northwestern University's Kellogg School of Management suggests that developments in AI research and the social sciences have not kept pace with one another.³⁰ Wang's study implies that AI-specific researchers are increasingly publishing their work within topic-specific forums while AI research is notably absent from references made by social scientists in their own work. This suggests that AI research is becoming isolated from the sociologists, economists, and philosophers who can best inform and benefit from developments within the field. Wang cautions that this growing gap will hinder the development of AI-driven technologies, arguing for increased collaborations between AI researchers and the broader social sciences that create a

²⁶ Noble, Safiya Umoja. *Algorithms of Oppression: How Search Engines Reinforce Racism*. New York: NYU Press (2018).

²⁷ Eubanks, Virginia. *Automating Inequality: How High-Tech Tools Profile, Police, and Punish the Poor*. New York: St. Martin's Press (2018).

²⁸ Obermeyer, Ziad, Brian Powers, Christine Vogeli, and Sendhil Mullainathan. "Dissecting racial bias in an algorithm used to manage the health of populations." *Science* 366, no. 6464 (2019): 447-453. <https://science.sciencemag.org/content/366/6464/447>

²⁹ Buolamwini, Joy, and Timnit Gebru. "Gender shades: Intersectional accuracy disparities in commercial gender classification." In *Conference on Fairness, Accountability and Transparency*, pp. 77-91. 2018. <http://proceedings.mlr.press/v81/buolamwini18a/buolamwini18a.pdf>

³⁰ Frank, Morgan R., Dashun Wang, Manuel Cebrian, and Iyad Rahwan. 2019. "The Evolution of Citation Graphs in Artificial Intelligence Research." *Nature Machine Intelligence* 1: 79–85.

two-way street of exchange. Other recent work, notably by Tim Miller, demonstrates ways in which the social sciences prove uniquely capable of sharpening and improving the capacity of AI technologies while refining a proactive analysis of their ethical dimensions.³¹

3. **Investments in areas of exchange and overlap between AI and the social sciences would support the broader goals of the European Commission surrounding innovation and oversight.** With the European Commission increasing its annual investments in AI by 70% through the Horizon 2020 research and innovation program, there is increasing support for AI research centers across Europe.³² As part of this effort, a corresponding investment is needed to enable existing social science research centers to develop research agendas that productively engage developments in AI. At the level of academic research, the social sciences are well-positioned to provide the insights and scrutiny needed to refine the ethical application of AI within society. Just as the AI-driven technologies can benefit from a clearer understanding of social, legal, and ethical challenges, AI applications can greatly enhance conventional social scientific research at a time when such research at institutions of higher learning across Europe is under increased financial strain and administrative scrutiny. Increased investments in applied and basic research related to AI applications across the social sciences would both improve the technologies under development. At the same time, industry-independent research provides credible evidence for oversight of tech companies.³³
4. **Building upon existing capacity within Europe would enable the creation of a public-facing, policy-focused research environment at the forefront of developments within AI-driven technologies.** Europe already possesses a robust independent technology assessment network driven by social scientists and higher education institutions in the form of the European Technology Assessment Group (ETAG), with the Institute for Technology Assessment and Systems Analysis (ITAS) at the Karlsruhe Institute of Technology (KIT) serving as its leading partner.³⁴ The ETAG provides scientific advisory services for the European Parliament across a broad spectrum of technology policy issues. Increased support should be directed toward ETAG to explicitly advance projects in partnership with European universities and research institutes that adopt a dual directional approach to the relationship between AI and the social sciences. ETAG is just one example of an existing research group dedicated to technology assessment with specific capacity in attending to issues surrounding AI. Other networks might also be considered for this function, particularly drawing upon the capacity of European research centers and universities. Beyond applied research for the purposes of creating specific technology assessment products

³¹ Miller, Tim. 2019. "Explanation in Artificial Intelligence: Insights from the Social Sciences." *Artificial Intelligence* 276: 1-38.

³² <https://ec.europa.eu/digital-single-market/en/artificial-intelligence>

³³ Matias, Nathan. "Why We Need Industry-Independent Research on Tech & Society." Technical Report, Citizens and Technology (CAT) Lab, Cornell University (January 2020).

<https://citizensandtech.org/2020/01/industry-independent-research/>

³⁴ <https://www.itas.kit.edu/english/etag.php>

and investigating the impacts of AI technologies upon society, dedicated support is also needed for basic research into areas of overlap between AI and the social sciences. Such investment in research on the governance and accountability of algorithms should be one of the priorities of the EU's next multiannual financial framework.

Conclusion

The Governance of AI Research Group would like to applaud the EU Commission's efforts in the direction of publishing The AI White Paper and providing a call for public consultation. As a group of scholars we are deeply concerned about the use of AI systems across many domains and governance systems. We see the power of AI applications and hope to work collaboratively with governments, companies, nonprofits, and other scholars to govern AI applications in such a way that the applications serve human interests and operate in a manner that protects human rights.

It is to this end, and in conjunction with what we perceive as overlap in our values with those of the EU Commission, that we have offered 4 general points for the Commission to consider as it seeks to improve the regulatory framework and governance structure around AI applications and systems. As the reader can see, we echo and detail two general points from the AI White Paper: (1) Increase AI Literacy Skills Within the Public Sector, the Private Sector, and the General Public to Enable Robust Auditing and Appeals Capability, and (2) Increase social science as this funding plays a critical role in meeting societal needs related to the development of general knowledge about AI applications and how they should be managed. We express our strong support for these cornerstones of effective implementation of a new regulatory framework.

While increases in human capital and funding of basic social science research help to improve upon the current AI governance system, strong, independent, and robust public agencies will also be required for the effective implementation of auditing algorithms and algorithmic decision making. In this section we argued that governance needs to take place across the inputs to the algorithm, the algorithm itself, and the use of the algorithm by individuals and institutions throughout society. This section highlights specific principles that should be followed and proposes institutions and the human capital needs to fill those institutions for implementation and enforcement of regulatory changes. Carefully considering both the governance structure of algorithms and the needed governance institutions are crucial aspects to consider in making this process successful.

Finally, we offer edits and improvements to the risk-based framework. While the members of the group generally endorse a risk-based approach, we note a number of deficiencies and provide numerous suggestions to make the framework more effective and overall reflective of both the opportunities and threats that AI applications present to society.
