Consumer Technology Association™

1919 S. Eads St.
Arlington, VA 22202
703-907-7600
CTA.tech

June 12, 2020

Consumer Technology Association Comments on
European Commission White Paper on Artificial Intelligence – A European Approach

The Consumer Technology Association® ("CTA")®[1] respectfully submits these comments in response to the European Commission ("Commission") request for comments and consultation on the White Paper on Artificial Intelligence – A European Approach ("White Paper").[2] CTA applauds the Commission's thoughtful framework set out in the White Paper. While detailed and attentive to important considerations, the framework should be recalibrated to reflect a more graduated approach which relies upon voluntary or consensus-based standards, voluntary governance and risk assessment processes, and a more nuanced framework that explicitly balances costs of over-regulating against the benefits of AI-based market solutions.

Overview and Summary

The Commission has undertaken important work and must be recognized for its initiative and efforts to enhance investment and development of artificial intelligence ("AI") applications and systems, while also working to increase public trust. Adoption of a formal framework can, if approached wisely, encourage greater innovation, enhance public trust and ensure responsible use. For that reason, CTA supports the Commission's stated intention "that the private sector should be fully involved in setting the research and innovation agenda" in order to provide the necessary level of co-investment.[3] However, these objectives can only be met by adopting a risk-based approach that is proportionate –i.e., that which *balances* potential harms against the many social and economic benefits created by AI.

The many benefits conferred by AI are evident. During the global COVID-19 pandemic, AI is driving important research and testing necessary to defeat the virus. For example, French AI company Iktos has partnered with SRI International, based in Menlo Park, to discover and develop new anti-viral therapies using deep-learning models,[4] and healthcare providers in New

[1] CTA® is the tech sector. Our members are the world's leading innovators—from startups to global brands—helping support millions of jobs. CTA owns and produces CES®—the largest, most influential tech event on the planet.
[2] White Paper on Artificial Intelligence, A European Approach to Excellence and Trust; COM(2020) 65 Final, Brussels 2.19.2020 ("*White Paper*").
[3] *Id*. at 7.
[4] *Iktos and SRI International Announce Collaboration to Combine Artificial Intelligence and Novel Automated Discovery Platform for Accelerated Development Of New Anti-Viral Therapies*, Drug Discovery Online (Mar. 4, 2020).

York have developed AI algorithms that can predict whether a COVID-19 patient is likely to suffer adverse events in the near future.[5]

The benefits of this technology extend well beyond healthcare. AI is now integrated in many parts of our day-to-day lives and is helping individuals and institutions increase productivity, efficiency, and outputs. The technology is also deployed across numerous different industries to improve healthcare and diagnostics, increase cybersecurity, reduce human trafficking, improve environmental sustainability, expand accessible services to persons with disabilities and automate dangerous tasks such as bomb-clearing or disaster relief. It is no wonder that the consulting firm McKinsey estimates that if Europe develops AI according to its current assets and digital position relative to the world, it could add some €2.7 trillion, or 20 percent, to its combined economic output by 2030.[6]

The exponential growth and reach of AI is due, in part, to the fact that the technology has developed in jurisdictions embracing the principle of "permission-less" innovation. Indeed, in the North American market, adoption of that principle in the United States has provided an operational environment that has permitted CTA members to deliver technology solutions that have changed the lives of millions of consumers. As CTA President Gary Shapiro has explained, this framework has proved to be a "winning formula for innovation in the United States and abroad."[7] The many benefits born of this environment offer an important reminder to the Commission that the need for any potential mandates or limitations must be weighed against the many benefits conferred by AI. For these reasons government policymakers must "think strategically about creating a regulatory environment that will encourage innovation in AI to thrive, continuing to prioritize cutting-edge research while taking into account disruptions AI could cause."[8]

At the same time, policymakers in the European Union and elsewhere should strenuously avoid designing AI regulations to achieve an objective of "technological sovereignty," as some in the European Commission have suggested. The United States and the European Union have the opportunity to set new global standards on AI governance through upcoming actions. It would be a grave mistake from an innovation and trade perspective to instead use new AI regulations as a vehicle to protect local champions and restrict market access by non-European providers of AI technologies and applications. Such efforts are unlikely to meaningfully strengthen the Europe's digital foothold and may stymie innovation in the very sectors the Commission seeks to promote.

To meet the goal of expanding beneficial uses while also enhancing trust and responsible use, the Commission should focus on: 1) supporting research/innovation (e.g., lighthouse center); 2) ensuring skills development and supporting SMEs; 3) promoting public sector adoption of this technology. In addition, the Commission should consider certain modifications to its framework including narrowing the scope of the framework; recalibrating to emphasize the use of a nuanced cost-benefit framework; relying upon industry consensus and voluntary standards; rejecting a

---

[5] *Coronavirus Tests The Value Of Artificial Intelligence In Medicine*, Fierce BioTech (22 May 2020).

[6] *Tackling Europe's Gap In Digital and AI*, Discussion Paper, McKinsey & Company (2019).

[7] Gary Shapiro, NINJA FUTURE at 91-92 (2018).

[8] *Id.* at 92.

strict liability standard for AI; and ensuring that AI already subject to existing industry regulations are not subject to duplicative or conflicting rules.

I. Scope and Definitions: The Commission Should Consider Narrowing the Scope of the Framework While Also Refining Key Definitions

As noted above, the breadth and scope of AI products and services is well documented. The broad range of applications and use cases available today means that AI is ill-suited for a top-down or "one-size-fits-all" governance framework. CTA agrees with the Commission's stated intent that "the new regulatory framework for AI should be effective to achieve its objectives while not being excessively prescriptive so that it could create a disproportionate burden." Overly burdensome regulatory frameworks that seek to mitigate or eliminate every potential risk are likely to hamper the development of useful AI technologies –innovations that offer substantial social benefits, including curing the sick, preventing deaths, and raising the standard of living for millions. The White Paper represents a necessary first step, but only addresses the challenges associated with AI at a very abstract and general level. CTA is concerned that these broad, general pronouncements in the White Paper are vague and may be difficult to implement. Accordingly, the Commission should adopt a more nuanced approach that recognizes that the breadth of AI use cases will require a focused response, which may differ by application.

Such an approach would embrace an application-specific framework for consideration of any new mandates. This approach would also ensure that AI policy accords with existing policy frameworks for the fields in which AI technologies are already in use today (such as healthcare, financial services or energy); that policymakers have sufficient expertise; and that the risk management calculations are tailored to the use case and the context-specific costs and benefits. Further, overarching rules that may apply to industries already covered by existing rules or principles (such as autonomous vehicles or online privacy) could conflict with or undermine existing frameworks. Application-specific approaches are more likely to be narrowly focused and tailored to the unique issues presented by AI in that use case.

The Commission's approach should focus on the challenges that are *unique* to AI and work to identify gaps not covered by existing regulation. Indeed, the EU has a wide body of existing law to cover general issues of consumer protection, privacy, and data security which already address many of the issues raised in the White Paper. A more effective use of the Commission's regulatory authority would be to identify and specifically address the novel issues that AI presents. For example, such efforts could explore methods to counteract algorithmic bias, or outputs that are inherently or systematically prejudiced due to erroneous assumptions or omission in an application's machine learning process. The Commission could also tackle issues of transparency by setting a framework for what constitutes explainable AI and the steps necessary to synthesize a machine learning process into a digestible concept which a consumer can comprehend to understand how a given result was achieved.

A. Recalibrate Risk-based Framework to Emphasize Import of Cost-Benefit Analyses

CTA reiterates its support for the Commission decision to use a risk-based approach to any potential new AI regulation, and encourages further dialogue and substantive consultations with the private sector on these questions. However, the framework articulated in the White

Paper should be recalibrated to ensure that any potential regulation is targeted to specific use cases, provides greater certainty, and does not discourage the development of AI. To accomplish these goals the Commission should consider taking the following steps.

1. *Evaluate the opportunity cost of not using AI*

Potential costs or harms associated with AI offerings must be weighed against such offerings' benefits. Such cost/benefit analysis is ubiquitous in policymaking across a wide range of fields whether or not it is always made explicit.[9] For example, in developing environmental policies policymakers weigh the burden on the environment against the cost of not producing many of the consumable goods essential to everyday life at a scale necessary to meet public need. The Commission must take on a similar analysis with respect to its AI policy.

2. *Utilize the principle of proportionality*

CTA commends the Commission for recognizing that not all AI applications require the same level of scrutiny. The probability of harm must be measured against the potential severity of harm, and the wider operational context when assessing risk. Utilizing that kind of proportional response will lead regulators to conclude that in many situations lower probability of harm may justify more limited regulation. The Commission should be sure to consistently apply this principle of proportionality throughout any new framework.

3. *Remove open-ended statements*

For example, the "exceptional instances" carve out[10] is unworkable and creates legal uncertainty because the concept is not tethered to specific use cases or applications. Similarly, the concept of "immaterial damages" in the risk definition[11] is also problematic because the standard is not used generally, and should be eliminated.

4. *Maintain reasonable performance standards*

Overly burdensome performance standards could discourage deployment of useful or necessary AI systems. Although AI systems have the potential to create unparalleled efficiencies and increase the likelihood of successful outcomes, the Commission should be careful not to measure AI by a performance standard that is perfection or its analog. Instead, the Commission must consider adopting a standard that holds AI system to perform with a level of accuracy equivalent to human performance, which leaves room for a reasonable margin of error. Standards should measure performance by evaluating whether AI is improving processes or scenarios. In other words, performance standards should serve to access whether the AI system is leading to outcomes that are better than the baseline human performance. These standards would be relative, rather than finite. Comprehensive and substantive consultations with the private sector is key to ensure regulatory intervention does not create a burdensome and ineffective regime, including a thorough examination of the type of requirements and enforcement mechanisms.

[9] Further, risk assessments must be holistic, and account for the potential that excess regulation may prevent certain AI-enabled products from getting to market, which have other unique social costs.

[10] *White Paper* at 18.

[11] *Id*. at 17.

5. *Implement consensus-based voluntary standards*

Where appropriate, the Commission should look to commonly accepted, or industry-developed, standards or best practices to govern AI obligations. Voluntary consensus-based standards can reduce the burden of complying with ill-formed regulation, eliminate the administrative costs of developing state-mandated standards, and decrease the overall cost of goods procured and the burden of complying with agency regulation. Indeed, CTA and its members have already been centrally involved in AI standard-setting, both in North America[12] and internationally.[13] These consensus-based standards often have broad support from industry, are more likely to reflect the most current technological developments, and reflect the most practical solutions available to the marketplace.

B. Refine Key Definitions

Another crucial consideration is to ensure that any future framework does not unintentionally reach traditional software, algorithms or other digital systems/processes that are not unique to AI. To do so, the Commission should consider adopting a definition of AI that narrowly focuses on those unique aspects of this technology (i.e., systems that are capable of learning on their own), but which avoids capturing general concepts or constructs used in computer science (i.e., algorithms) but which do not, on their own, constitute AI. To that end, CTA urges the Commission to consider utilizing the definition used in its April 2018 Communication on AI:

> *Artificial intelligence (AI) refers to systems that display intelligent behaviour by analysing their environment and taking actions – with some degree of autonomy – to achieve specific goals. AI- based systems can be purely software-based, acting in the virtual world (e.g. voice assistants, image analysis software, search engines, speech and face recognition systems) or AI can be embedded in hardware devices (e.g. advanced robots, autonomous cars, drones or Internet of Things applications).*[14]

This definition is both broad enough to encompass the full spectrum of AI systems without capturing established technologies which are already subject to existing legislation. At the same time, this definition avoids the use of certain terms that were utilized in the High-Level Expert Group on AI's definition,[15] such as "software [] and … hardware systems" which may unintentionally capture systems that do not involve AI.

---

[12] For example, CTA is developing standards focused on AI in healthcare. *See* The Use of Artificial Intelligence in Health Care: Trustworthiness; or Definitions and Characteristics of Artificial Intelligence; *and*, Riya Anandwala and Danielle Cassagnol, CTA, Press Release, CTA Launches First-Ever Industry-Led Standard for AI in Health Care (rel. Feb. 25, 2020), available at https://www.cta.tech/Resources/Newsroom/Media-Releases/2020/February/CTA-Launches-First-Ever-Industry-Led-Standard.

[13] For example, CTA has been actively participating in ISO/IEC JTC 1/SC 42, the international standards committee responsible for standardization in the area of Artificial Intelligence.

[14] COM(2018) 237 final, p. 1 (cited in *White Paper* at n. 46).

[15] "Artificial intelligence (AI) systems are software (and possibly also hardware) systems designed by humans that, given a complex goal, act in the physical or digital dimension by perceiving their environment through data acquisition, interpreting the collected structured or unstructured data, reasoning on the knowledge, or processing the information, derived from this data and deciding the best action(s) to take to achieve the given goal. AI systems can either use symbolic rules or learn a numeric model, and they can also adapt their behaviour by analysing how the environment is affected by their previous actions." (cited in *White Paper* at n. 47).

Other key concepts in the White Paper would similarly benefit from further development. Concepts such as "human agency and oversight," "robustness and accuracy," "transparency," "diversity," "non-discrimination" and "fairness" are core concepts, yet very few are formally defined. Further consideration must be given to how these terms are defined in the context of a balanced, risk-based framework. CTA is keen to see and provide further feedback on the specific language at legislative and regulatory stages.

An effective framework would also clarify the degree to which these principles should be present in a particular use case, recognizing that not all applications require the same degree of scrutiny. The Commission should leverage the existing work of CTA and other entities active in developing consensus-based standards, protocols and best practices for AI, many of which include workable definitions of these concepts. For example, CTA's recently published standards document *Definitions/Characteristics of Artificial Intelligence in Health Care*[16] defines several of the terms noted above in a workable, flexible manner that may assist the Commission in further refining its framework.

II. Mandates Under the Framework: Certain Elements of the Framework Are Overbroad and May be Unnecessary

As the Commission aptly notes, any regulatory intervention imposed on relevant actors must be "focused and proportionate." The applications of AI systems span across numerous different industries –many of which are already subject to existing legislation, regulation, and/or agreed upon industry standards. The Commission must be careful not to implement a framework that is too broad or it will conflict with existing law.

A. Keeping of Records and Data

The White Paper proposal to require retention of datasets would conflict with certain GDPR provisions that require the deletion of personal data. Such requirements may also interfere with the use of copyrighted datasets, which often only permit short-term access to the data at issue. If the Commission were to impose a framework that effectively forces data to be collected and stored centrally, it runs the risk of eliminating privacy benefits built-in to the AI system, such as on-device processing. Such a framework would also prevent the use of off-the-shelf, open-source models, since developers would generally have no access to the data used to train them. In lieu of extensive data and records retention mandates, the Commission should consider the use of simple notice or transparency duties to address the need for a certain level of explainability. For example, an AI developer may be required to provide certain disclosures or transparency notices, and AI deployers further up the service delivery chain would then need to ensure that those disclosures flow through to the public and that their terms and conditions are consistent with such disclosures. Commission guidance on implementing such an approach in the context of a service delivery chain, including identifying potential roles of the respective parties in such supply chain, would be useful.

[16] *Definitions/Characteristics of Artificial Intelligence in Health Care*, ANSI/CTA Standard, ANSI-CTA-2089-1, available at: https://shop.cta.tech/products/definitions-characteristics-of-ai-in-health-care.

B. Training Data

Several elements of this proposal remain impractical in light of restrictions under current EU legislation. This includes the proposed obligation for developers to "ensure datasets are sufficiently representative." This proposal may conflict with GDPR rules that prohibit developers from accessing sensitive personally identifiable information and attributes, such as gender, ethnicity or income levels.[17] Rather than putting requirements on training data, it would be better to base requirements on testing model performance using benchmark datasets. This would ensure that the outputs are within an acceptable range, since it is the model output that ultimately determines the real world impact of an AI system. CTA looks forward to further development of the Commission's intentions to address undue restrictions on data flow through bilateral negotiations or WTO actions. Access to data is critical to continued development of AI systems, and also helps to ensure that data used for training is representative and diverse.

C. Robustness and Accuracy

CTA shares the Commission's goal to ensure trustworthy AI systems. Any regulatory framework must be sufficiently flexible to reflect how these systems are developed. For example, adoption of blanket "requirements ensuring that outcomes are reproducible" would be unreasonable in light of how certain AI systems are developed. Many systems have randomness built in, which makes it impossible to guarantee an identical output every time even if the input is the same. The Commission should avoid rigid requirements that would impose an expectation for all model outputs to be reproducible, and instead create a workable framework that considers how a given systems is developed.

D. Human Oversight

As the White Paper correctly notes, not all AI systems pose similar risk. Accordingly, any requirement that AI systems have human oversight should be reserved for certain systems in the high-risk category where human oversight would not increase risk or obviate the technology's purpose altogether. Rather than adopting a broad mandate for human oversight in all settings, the Commission should consider defining standards for the development of AI systems and the assurance of reliability. Then the Commission must carefully consider, and define, the necessary role and extent of human oversight in a specific context. Such obligations must be practical and appropriate, and only mandated when absolutely necessary. For example, human oversight may not be necessary (or appropriate) for AI systems that make content recommendations or deliver routine/non-consequential services or applications. However, human oversight may be necessary for high-risk applications, such as robotic surgery. Further, any obligation to maintain human oversight and/or intervention should apply during the *training* of fully autonomous systems, but not necessarily during the period of commercial operations (when the systems have already been validated for safe operation).

[17] Further, the Commission must provide clarity on what constitutes "sufficient" representation. The answer will likely vary based on different use cases and applications.

III. Liability Issues: Adoption of a Strict Liability Framework Is Not Necessary and Would Likely Suppress Investment and Development of AI

A. The Commission Should Drop Any Consideration of a Strict Liability Framework for AI Systems

The White Paper suggests that the Commission may need to consider adopting a strict liability framework for AI systems.[18] But doing so would be unwise because current liability systems are sufficient to allocate responsibility and liability among AI developers and deployers. Further, adoption of a strict liability framework would likely undermine other Commission goals to support further private and public sector investment in trustworthy AI.

Adoption of a strict liability rule for AI systems is unnecessary because the current liability framework in the EU is successful and works well across a range of technology sectors. The current framework properly assigns liability to economic actors in a manner that ensures both accountability by economic actors and recourse for affected individuals. At the same time, the current system provides certainty and clarity for economic actors that fail to meet those standards. In addition, the current framework is a technology neutral framework that does not favor one form of technology over another. Further, there is no evidence that the current framework does not properly allocate responsibility, or that it would not do so when applied to AI systems.

While the White Paper suggests that certain AI systems may lack the necessary transparency to assign liability, this point fails to account for the fact that developers of AI systems will have incentives to put in place contractual commitments to one another. Indeed, the White Paper recognizes that many actors are involved in the lifecycle of an AI system including entities: *developing* the technology, those *deploying* the technology and potentially *others*, such as distributors, resellers, service providers and individuals.[19] Under this framework private economic actors in the development chain will have allocated responsibility for any system failures amongst themselves. That, in turn, will ensure that consumers still have recourse against private economic actors involved in the development of such systems. Either the entity providing the end product or service to consumers will be liable, or that liability will rest with an entity developing the AI system, component or feature. Under either scenario, the roles and responsibilities of stakeholders in the AI ecosystem can allocate, by contract, liability among the developers and deployers of the AI.[20] For these reasons, while some changes to existing liability regulation for AI may be necessary to reflect the unique challenges posed by AI systems and its components, such changes need not extend to adoption of strict liability principles.

---

[18] *White Paper* at 15 ("… the Commission is seeking views whether and to what extent it may be needed to mitigate the consequences of complexity by adapting the burden of proof required by national liability rules for damage caused the by the operation of AI applications.").

[19] *Id*. at 22.

[20] This framework also eliminates the potential, which could arise in a strict liability framework, that even smaller entities who have only limited involvement in developing an AI system could be liable for outcomes for which they had no knowledge or control. The same principle counsels against attempts to revise the EU Product Liability Directive to address issues unique to AI.

Neither the White Paper nor the accompanying Report on the Safety and Liability Implications of Artificial Intelligence, the Internet of Things and Robotics,[21] evaluates the potential ramifications of adopting a strict liability framework for all AI systems. Such a decision could have wide-ranging implications that could seriously undermine the development and expansion of important AI systems. Strict liability frameworks are most often utilized in situations where the potential harm to individuals is significantly hazardous.[22] In such circumstances questions of intent or negligence may, understandably, be less meaningful. But AI systems do not categorically create the kinds of hazards used to justify strict liability standards in other systems. Instead, the broad range of AI applications and uses cases, some of which could have ill effects but others not, is not well suited for strict liability principles.

Finally, the Commission must clearly delineate safety and liability concepts from other important, yet distinct concerns. For example, the White Paper couples non-liability issues, such as ethics, privacy, cyber security and mental health, with liability issues.[23] However, these issues can and should be addressed separately, outside of the framework for assigning liability duties.

B. Benchmarks for AI Performance Should Be Reasonable and Avoid Overly Demanding Standards

Evaluating the performance and trustworthiness of AI systems requires adoption of certain standards or metrics. The Commission's consideration of such issues is appropriate, but any benchmark for measuring AI system performance should not be perfection. Instead, reasonable and achievable improvements on current levels of human performance should be the goal. AI can *improve* over human decisions in many areas, but requiring that such systems never make a mistake will deter new research, investment and development.

IV. Enforcement Issues: Commission Must Reconsider *Ex Ante* Conformity Assessments In Lieu of Voluntary Self Assessments

A. Mandatory *Ex Ante* Conformity Assessments Are Not Aligned with a Risk-Based Approach

The White Paper suggests that "prior conformity assessments" are necessary to "verify and ensure that certain mandatory requirements applicable to high-risk applications are complied with."[24] However, mandating potentially overly prescriptive *ex ante* conformity assessments would be costly, imprecise and potentially undermine the exponential growth and reach of current AI systems and applications. Accordingly, the new regulatory framework for AI should be effective to achieve its objectives, without being excessively prescriptive in a manner that could create a disproportionate burden.

In lieu of mandated *ex ante* conformity assessments the Commission should consider alternative frameworks in favor of more balanced risk-based approaches that adequately address costs and benefits of new rules. For example, a modified framework that embraces voluntary *ex ante* self-assessments for risk, coupled with an *ex post* enforcement framework for the highest

---

[21] COM(2020)64, Report on the Safety and Liability Implications Of Artificial Intelligence, The Internet of Things and Robotics (rel. 19 Feb. 2020).

[22] *See, e.g.*, EU Product Liability Directive.

[23] *White Paper* at 15.

[24] *White Paper* at 23.

risk AI applications, would likely achieve similar results without undermining innovation and growth in this field. Such a framework would incentivize the adoption of voluntary risk-assessments and governance programs by AI developers and deployers. These entities would have strong incentives to conduct rigorous self-assessments and implement strong governance programs in advance of AI products and services reaching the market in order to avoid potential *ex post* enforcement by the Commission after these AI products reach the market.

This approach would have the added benefit of employing a flexible and contextually-specific risk assessment process to the particular AI-enabled product or service going to market.[25] In contrast, under the framework articulated in the White Paper, there remains the potential for *ex ante* risk assessments that would lack necessary specificity and could result in a one-size-fits-all assessment that does not properly assess risks or benefits of the specific system.

Alternatively, the Commission should at the very least, limit *ex ante* conformity assessments to AI applications deemed to be very high-risk and which are also based on specific articulated principles (analogous to the requirement for data protection impact assessments under GDPR). However, even this approach is problematic, in part because there is considerable uncertainty in relation to the anticipated categorization of AI applications as high-risk. As noted, the risk profile between content recommendation systems and robotic surgery are quite different.

B. Conformity Assessments Should Account for Existing Products and Those Under Development

The White Paper does not account for products already in the market. Presumably the Commission intends for its framework to apply prospectively, in part because extending obligations to existing services and products already in the market or under development would create significant operational costs for developers and deployers that may have to retrofit existing AI systems to adhere to new standards. Further, the potential for a logjam of retrospective testing and risk assessments is quite possible if the framework was not clearly applied on a proactive basis. It is therefore important that the Commission makes clear that any new obligations arising out of this initiative would apply only prospectively, and that all existing AI-based services and products are expressly excluded via a grandfather clause.

Further, the Commission should expressly exclude any AI-based services and products that are in the pre-market phase. Thus, any AI that is currently subject to research, development and pre-market product development should be expressly excluded from any obligations. This exclusion will affirm the important principle that developers and deployers can conduct confidential testing and pilot projects of AI applications and service prior to any obligation to perform conformity assessments.

[25] A contextually-specific risk assessment would also be consistent with the Commission's finding that any new framework must not be overly prescriptive and create a "disproportionate burden" on affected parties. *White Paper* at 17. Consistent with that approach, risk assessments that are contextually-specific should provide a holistic assessment of the "potential harms, the magnitude of those harms, the technical state of the art, and the potential benefits" associated with the specific application of AI that is being evaluated.

C. Retraining on European Datasets is Problematic

Finally, the White Paper seems to suggest that AI systems which fail conformity assessments may need to be retrained in Europe raises significant concerns.[26] This proposal should be rejected in its totality.

There is no guarantee that European datasets or training that takes place on European soil would provide any measurable improvements to the performance of an AI system. Indeed, to the contrary, precluding the use of foundational data sets that include data and information derived from outside of Europe would risk reducing a system's performance by unduly narrowing or restricting available training data. It is axiomatic that AI and the machine learning systems that support such technology are more accurate, precise and equitable when trained on larger, rather than smaller data sets. Thus, requiring retraining on data derived only in Europe could actually exacerbate the risk of bias, discrimination or inaccurate outcomes and thereby reduce the quality and efficacy of the overall AI system. That result is not in the interest of the Commission or EU residents.

V. Other Issues: Avoid Duplication of Existing Mandates Already Applicable to AI

A. Exclude from Scope of Any New Rules Those Applications Already Subject to Oversight and Established Governance Frameworks

Given the scope of AI applications and use cases across numerous industries, the potential for conflicting or duplicative laws in certain markets is problematic. To address this concern the Commission should carefully review and identify those areas, such as autonomous vehicles, which may already have existing duties and obligations. For example, the creation of conformity assessment procedures for high-risk AI applications, if applied to autonomous vehicles, could be duplicative of other existing requirements applicable to autonomous vehicles. Accordingly, any general rules adopted by the Commission should carve out certain areas, such as autonomous vehicles, which are subject to other regulatory requirements.

Indeed, any regulation of AI used in autonomous vehicles as a separate framework would conflict with the existing legal framework for such applications. Specifically, the "type approval" regulatory framework, which is based on internationally harmonized rules governs the safety and operation of both manual and automated vehicles. That framework should continue to apply without interference, or conflict from any new mandates the Commission may consider.

B. Mandating Ongoing Human Oversight of Fully Automated Vehicles Is Inappropriate and Unnecessary

The White Paper also suggests that the Commission should adopt human oversight obligations to autonomous vehicles.[27] As noted above, utilizing the element of human oversight in certain circumstances may be appropriate. However, applying that principle to autonomous vehicle systems is fundamentally at odds with the objective of developing and deploying fully automated vehicles. Notably, the EU General Safety Regulation distinguishes "automated vehicles" from "fully automated vehicles" and makes it clear that fully automated vehicles are

[26] *White Paper* at 23.

[27] Suggesting that human oversight would be mandated in the "monitoring of the AI system while in operation and the ability to intervene in real time and deactivate (e.g. a stop button or procedure is available in a driverless car when a human determines that car operation is not safe)." *White Paper* at 21.

"designed and constructed to move autonomously without any driver supervision."[28] Accordingly, CTA recommends that any considerations of human oversight and/or intervention for AI should expressly exclude any fully automated systems which are covered by existing rules. Human oversight may be appropriate during the training of these systems but would not be appropriate during deployment in commercial situations after the systems have already been validated for safe operation.

VI. Conclusion

CTA and its members have a significant interest in ensuring that European consumers benefit from AI-powered products and services. The Commission should proceed carefully to ensure that its policies promote continued development and deployment of AI products and services that enhance the lives, safety and interests of European consumers. CTA stands ready to continue its central role in the development of consensus-based standards that advance these goals and promote continued dynamic growth and innovation throughout the consumer technology industry.

/s/ Douglas K. Johnson
Douglas K. Johnson
Vice President, Technology Policy

/s/ Michael Petricone
Michael Petricone
Sr. Vice President, Government and Regulatory Affairs

[28] Regulation (EU) 2019/2144 of the European Parliament and of the Council, On Type-Approval Requirements for Motor Vehicles and Their Trailers, and Systems, Components and Separate Technical Units Intended for Such Vehicles, as Regards Their General Safety and the Protection of Vehicle Occupants and Vulnerable Road Users at Art. 3, sub. (22) (rel. 27 Nov. 2019).